

How to Write an LLVM Backend #4: Instruction Selection

Nov 17, 2020

We previously provided a brief overview of the compilation process in the LLVM backend. In this post, we will take a closer look at the first stage of that process: instruction selection. The aim is to understand how this works and how we can configure it before we look at the concrete implementation of our RISCW backend.

NOTE: The LLVM documentation has this short and clear explanation of how instruction selection works, so this post is mostly a re-statement of that enhanced with examples.

NOTE: The examples shown in this post are constructed using our skeleton RISCW backend. The source for that can be found [here](#).

NOTE: Posts in this series:

1. Introduction
2. Getting Started
3. Setting Up a New Backend
4. Configuring the Build System
5. Instruction Selection
6. Arithmetic Instructions

The instruction selection process takes as an input the LLVM IR and outputs a sequence of instructions that use an infinite set of registers. The process is split into the following phases:

1. Build initial DAG
2. Optimize
3. Legalize types
4. Optimize
5. Legalize operations
6. Optimize
7. Target instruction selection
8. Scheduling and formation

I think it is easier to learn what's happening by following an example. So we will consider how the simple C program below is transformed as instruction selection proceeds. The code has a function MUL that takes in a 64-bit argument `x` and a 32-bit one `y`. The arguments are multiplied and a 32-bit integer result is returned.

```

unsigned int MUL(unsigned long long int x, unsigned int y)
{
    return x * y;
}

```

Build DAG Phase

This is the first phase of instruction selection. It takes an LLVM IR as an input and produces a Selection Directed Acyclic Graph (DAG) as an output. Every other phase in the instruction selection process operates on the DAG until the output instruction sequence is emitted. As discussed a few posts ago, the LLVM IR is generated from the C code by a front-end tool, like Clang, and later optimized by the LLVM optimizer. Here is the LLVM IR for our simple C program:

```

define dso_local i32 @MUL(i64 %x, i32 %y) local_unnamed_addr #0 {
entry:
    %0 = trunc i64 %x to i32
    %conv1 = mul i32 %0, %y
    ret i32 %conv1
}

```

The Selection DAG is effectively a fancy tree data structure that represents a basic block in the LLVM IR. A basic block is a sequence of instructions without branch destinations (except the entry) and no branch instructions (except the exit). Our example MUL function is simple enough that it only has one basic block called `entry`, but other functions typically have multiple basic blocks. For instance, the hello function below has four basic blocks: `entry`, `if.then`, `if.else` and `return`. Each basic block will be translated to a separate DAG.

```

define dso_local i32 @hello(i32 %x) local_unnamed_addr #0 {
entry:
    %cmp = icmp eq i32 %x, 100
    br i1 %cmp, label %if.then, label %if.else

if.then:                                ; preds = %entry
    %call = tail call i32 @bitcast (i32 (...)* @hello100 to i32 (i32*)) (i32 100)
    br label %return

if.else:                                ; preds = %entry
    %call1 = tail call i32 @bitcast (i32 (...)* @hello0ther to i32 (i32*)) (i32 %x)
    br label %return

return:                                  ; preds = %if.else, %if.then
    %retval.0 = phi i32 [ %call, %if.then ], [ %call1, %if.else ]
}

```

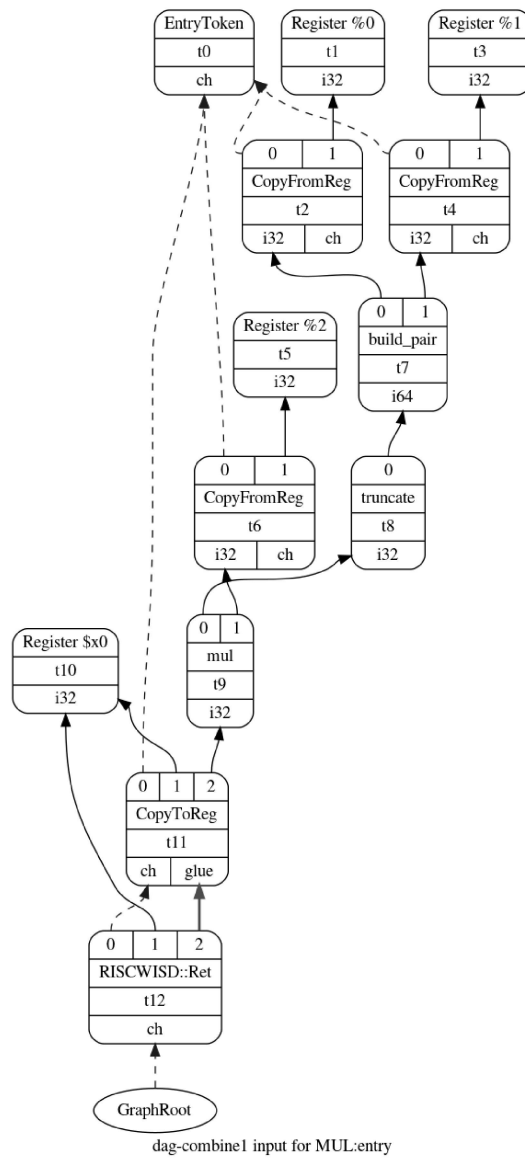
```
ret i32 %retval.0
}
```

A Selection DAG has the following properties:

- Each node is an instance of the `SDNode` class representing an operation, like add, subtract, multiply, etc. The operation types are in `include/llvm/CodeGen/ISD0pcodes.h`.
- Each node has 0 or more input operands, which are outgoing edges to the node defining the used value. Edges are instances of the `SDValue` class.
- Values produced by the operations in a node have a Machine Value Type (MVT). Examples of types are `i1` and `i8` which represent 1- and 8-bit integer values respectively.
- Nodes that have side-effects force an ordering on operations, e.g return statements and loads, so they have a special *chain* value of type `MVT::Other` as both an input operand and an output.
- A leaf node of type `ISD::EntryToken` is the designated entry into the code.
- The root node of the DAG is the final side-effecting node with a chain operand. This is the last operation performed in the block of code represented by the DAG, such as a return at the end of a function.

WARNING! The type system in the LLVM backend is *VERY* limited. A lot of useful type information is thrown away when the input LLVM IR is converted to a DAG. Perhaps the most notable omission are pointer types which are completely absent from the list of types in the `MVT` class. As a result, pointers have integer type in the DAG, so it is difficult to tell whether a node, like an add, has integer or pointer operands.

Here is the initial Selection DAG that I get from the MUL function.



The entry node is at the top of the image while the root of the DAG is at the bottom. There is also a chain of nodes joined with blue edges starting from the root and ending on the entry node. These blue edges are the chain operands discussed earlier. In contrast, the black edges show the flow of data values such as integers and floating-point numbers. An equivalent text representation of this Selection DAG is shown below:

```
t0: ch = EntryToken
    t2: i32,ch = CopyFromReg t0, Register:i32 %0
    t4: i32,ch = CopyFromReg t0, Register:i32 %1
    t7: i64 = build_pair t2, t4
    t8: i32 = truncate t7
    t6: i32,ch = CopyFromReg t0, Register:i32 %2
    t9: i32 = mul t8, t6
    t11: ch,glue = CopyToReg t0, Register:i32 $x0, t9
    t12: ch = RISCWISD::Ret t11, Register:i32 $x0, t11:1
```

NOTE: A Selection DAG may contain target-independent and target-specific operation nodes. Target-independent operations are defined in `ISD0pcodes.h` as mentioned above. Target-specific operations are defined by each backend, typically in the same file that

declares the target's TargetLowering subclass. You can find the target-specific nodes for RISCW in `llvm/lib/Target/RISCW/RISCWISelLowering.h`. The Selection DAG shown above for MUL has a node `t12` with the target-specific operation `RISCWISD::Ret` that represents the return from the function.

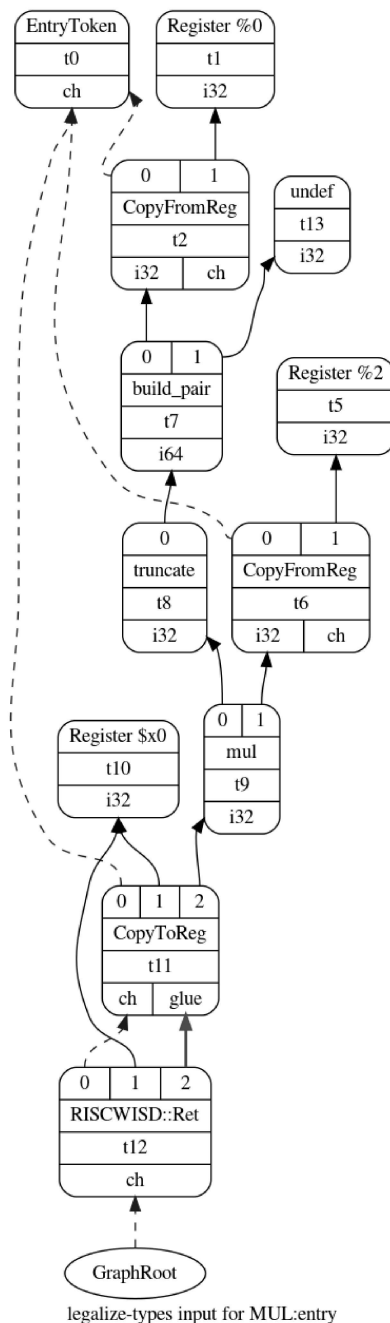
NOTE: You can tell LLVM to generate a visual representation of the Selection DAG at various stages of the instruction selection process by passing `-view-dag-combine1-dags`, `-view-legalize-dags`, `-view-dag-combine2-dags`, `-view-isel-dags` or `-view-sched-dags` as command line arguments when invoking `llc`. The `-debug` argument tells `llc` to generate a text representation of the Selection DAG.

Optimization Phase

There are three optimization phases during the instruction selection process. The first optimization pass is applied immediately after the Selection DAG is constructed from the LLVM IR. The remaining two passes are applied after the legalization phases. According to the LLVM documentation, these optimization phases are meant to simplify unnecessarily complex Selection DAGs that may have been produced by other phases, like legalization.

As far as I can tell, these optimization phases are limited to running a compilation pass that *combines* groups of nodes into simpler nodes. To achieve this, LLVM has this *huge* C++ file (>20,000 lines) that traverses the DAG and uses pattern matching on the nodes to identify and apply optimization opportunities. For example, the combiner could convert the nodes `(add (add x, y), z)` from the input Selection DAG into `(add x, y+z)` thus eliminated an `add` node.

The output Selection DAG of the first optimization pass is shown in the figure below for our example MUL function. Compared to the Selection DAG from the previous section, the optimized DAG contains the node `t13` instead of nodes `t3` and `t4`. The combiner was able to perform this optimization, which discards the most significant 32 bits of the 64-bit function argument `x`, because the multiplication in the function body and the return argument rely on 32-bit integers. Therefore, the DAG can be simplified by eliminating `t3` and `t4`.



legalize-types input for MUL:entry

LLVM's DAG combiner only optimizes target-independent operation nodes, i.e. those defined in `ISD0pcodes.h` like add, sub, load and store. Backends can provide additional combiner functionality for target-independent and target-specific nodes by overriding the `PerformDAGCombine` hook in their `TargetLowering` subclass. Also, the backend has to specify the target-independent nodes that it wishes to be notified of during an optimization phase by calling `setTargetDAGCombine` in the constructor of their `TargetLowering` subclass.

NOTE: Confusingly, the `TargetLowering` subclass of an LLVM backend is usually implemented in the `<TARGET_NAME>ISelLowering.cpp` file. You can find the RISCW one in `llvm/lib/Target/RISCW/RISCWISelLowering.cpp`.

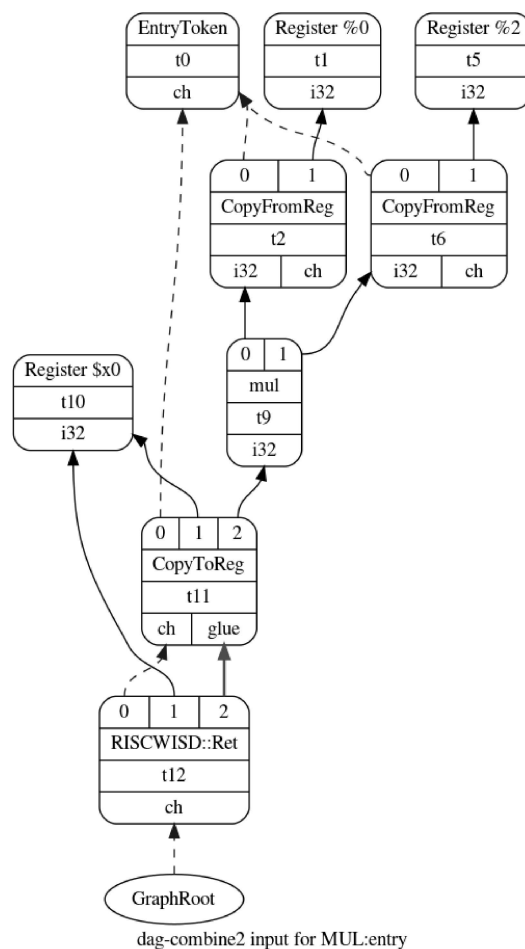
NOTE: The XCore backend has a good example –and by *good* I mean *easy-to-read*– for using `PerformDAGCombine` and `setTargetDAGCombine` (see here and here). Also, take a look here to see how that backend combines the nodes `add(add(mul(x, y), a), b)` into the simpler

target-specific `lmul(x, y, a, b)` node as the XCore architecture has a long multiply-accumulate instruction.¹

Legalize Types and Legalize Operations Phases

The optimized DAG shown above contains a node `t7` that produces a 64-bit value of type `i64`, but the RISCW machine is only 32-bit. The legalize phases take care of these problems.

The legalize types phase is run first. It converts the DAG to only use data types natively supported by the machine. To achieve this, it converts, or *promotes*, small types to larger ones; for example, a 1-bit integer of type `i1` is converted to a 32-bit integer of type `i32`. Also the compiler breaks up, or *expands*, large integers into smaller ones, such as 64-bit integers of type `i64` that are converted into 32-bit integers of type `i32` in a 32-bit machine. The DAG below is a legalized version, with regards to types, of the optimized DAG shown before for the MUL function. In this case, the compiler decided to eliminate the `t3` and `t7` nodes to ensure that the new DAG only uses `i32` integers.



The legalize operations phase is run after the second optimization phase. It converts the DAG to only use operations supported by the machine. For example, a DAG may contain bit-rotate left (`rotl`) nodes, but the target instruction set might not have that instruction. Thus, the operation needs to be implemented with a combination of bit shifts and ors; the legalize operations phase implements such changes to the DAG.


There are three strategies to legalize operations. First, the compiler can *expand* an unsupported operation into an alternative set of supported operations that emulate the required functionality. Second, a type can be *promoted* to a larger type that supports the missing operation. And third, the target can use a hook in its TargetLowering subclass to implement a *custom* legalization in C++.

NOTE: In future posts, I will explain how a backend can configure the legalization phases; this can be done in the constructor of the target's TargetLowering subclass.

Target Instruction Selection Phase

At this point, the DAG mostly contain nodes with target-independent operations, like add and subtract, along with a few target-specific operations, like `RISCWISD::Ret`. So we need to map these 'abstract' operations to concrete machine instructions from our target architecture; LLVM takes care of this in the instruction selection phase. The idea is fairly simple: the compiler pattern matches the nodes in the DAG to machine instructions. The patterns and descriptions of the instructions are provided by the developer of the compiler backend via code written in TableGen. Also, more complex pattern matching that is difficult to express using TableGen can be coded directly in C++.

We will look at TableGen more closely in future posts. For now, let us consider how LLVM matched the legalized and optimized DAG for MUL function to the new DAG shown below. In this case, only two nodes needed changes. `t9` had its `mul` operation replaced with the instruction `MUL` while `t12` had `RISCWISD::Ret` replaced with the instruction `PseudoRET`.



Scheduling and Formation Phase

Thankfully, there is not much to say here as the post was getting quite long! This phase turns the DAG into a list of instructions according to some target constraints. For example, the target can specify its scheduling preference by calling `setSchedulingPreference` from the constructor of its `TargetLowering` subclass.

WARNING! Make sure you experiment with various scheduling preferences! Choosing the wrong one can significantly degrade the emitted code if the scheduling strategy does not match the features of your architecture and processor. The various scheduling options are listed here.

Below is the emitted list of instructions for our MUL program. There are a couple of important things to note. First, the emitted code is still using an infinite set of virtual registers as opposed to actual machine registers; the compiler takes care of this problem later with register allocation. And second, there is already some liveness information that is useful for the register allocation process; for example, there are two registers, i.e. `x0` and `x2`, which are live at the start of the basic block.



```
bb. 0. entry:
  liveins: $x0, $x2
  %2:gpr = COPY $x2
  %0:gpr = COPY $x0
  %3:gpr = MUL %0:gpr, %2:gpr
  $x0 = COPY %3:gpr
  PseudoRET implicit $x0
```

Notes

1. Take a look at the XCore XS1 architecture reference manual for a detailed description of the multiply-accumulate (LMUL) instruction. [↗](#)

Source Code Artisan

Source Code Artisan
andresag1118@gmail.com

 andresag01
 andresag

A personal blog on computing topics